Unforced Navier–Stokes

Compact symmetric stresses, exact coefficient margins, and formal verification on the three-torus

Tony Newton
Newton Astro Labs, United Kingdom

Abstract

The absence of external forcing makes the three-dimensional Navier–Stokes regularity problem substantially more restrictive than a construction with prescribed smooth forcing. We study exact mathematical prerequisites for a candidate unforced finite-time singularity on the periodic three-torus. An explicit compact symmetric-divergence inverse is proved for smooth axisymmetric meridional sources supported in an annulus, with vanishing total axial force as the necessary and sufficient compatibility condition. A separate two-diagonal pressure-neutral construction exhibits a sharp radial sign obstruction. Finite seven-ray stress algebra yields the exact coefficient margin 15971/122900 and sharp componentwise amplification factor 14. A compact source-matched dual test excludes a strong physical entrywise stress bound at the corresponding numerical scale. For the frozen cylindrical seven-ray frame, an exact matrix reconstruction identifies that bound with a sufficient four-coordinate H8 stress estimate. We give the exact three-inequality description and sharp dual support function of the original seven-coefficient cone. An exact rational Bernstein certificate shows that even this larger absolute coefficient budget fails for the specified manufactured meridional source. This does not decide the authentic H8 residual. Lean additionally checks the meridional pointwise inverse and both integral antiderivative identities, alongside the finite cone algebra, physical residual expansion, and conditional terminal-speed retention. These results do not yet construct the source-matched, scale-sharp, positive-cone-compatible compact stress or a genuine zero-force singular solution. The global unforced existence question remains open.

1 The problem and the mathematical objective

We work at viscosity one on the unit three-torus $\mathbb{T}^3 = \mathbb{R}^3/\mathbb{Z}^3$. For a smooth velocity $u(t,x) \in \mathbb{R}^3$ and pressure $p(t,x) \in \mathbb{R}$, the unforced incompressible equations are

$$\partial_t u + (u \cdot \nabla)u - \Delta u + \nabla p = 0, \qquad \operatorname{div} u = 0. \tag{1}$$

The specific target is to find nonzero smooth divergence-free periodic initial data and a solution smooth for $0 \le t < 1$, of finite kinetic energy on a fundamental cell, whose speed becomes unbounded as $t \uparrow 1$. The source-free condition in (1) is essential: a verified construction with a nonzero smooth force does not establish this target. The Clay problem statement makes this distinction explicit [1]. The published OpenAI Lean formalization states Navier–Stokes breakdown with smooth forcing; its unforced finite-time result concerns Euler rather than Navier–Stokes [2].

Our approach isolates small results that are meaningful independently of the unresolved global theorem. First, a source must admit a compact symmetric-divergence correction. Second, if that correction must remain in a prescribed positive stress cone, its components need quantitative control relative to a density envelope. Third, all terms in the actual nonlinear physical residual must cancel with a single periodic pressure. Finally, a singular leading trace

can survive only if the correction is provably subleading. Sections 2–7 establish exact parts of this chain.

The construction of compact symmetric anti-divergence operators under force and torque compatibility is not new; see, for example, the support-preserving operator and its explicit moment correction in [3, §3.6 and Appendix A]. The direct axisymmetric formula below is offered as an elementary source-specific coordinate realization, not as a priority claim for compact symmetric anti-divergence. The new obstacle for the present programme is the additional shrinking-support, weighted-seven-ray positivity requirement, which is not supplied by general existence of a compact tensor. Section 5 also shows that a natural *stronger* physical-component envelope is incompatible with one explicit source, even modulo pressure gradients. We subsequently exclude the original absolute coefficient envelope for the same manufactured source by a sharper dual calculation; the actual H8 source remains a separate problem.

2 A compact symmetric inverse in an annular region

Let $0 < a < b$ and $c < d$, and write $\Omega = (a, b) \times (c, d)$ in meridional cylindrical coordinates (r, z). All functions in this section are independent of the azimuthal angle θ. The orthonormal cylindrical frame is (e_r, e_θ, e_z). For a symmetric tensor with components

$$S = K\, e_r \otimes e_r + L\, e_\theta \otimes e_\theta + M(e_r \otimes e_z + e_z \otimes e_r) + C\, e_z \otimes e_z, \tag{2}$$

standard cylindrical differentiation gives

$$(\operatorname{div} S)_r = \partial_r K + \frac{K - L}{r} + \partial_z M, \tag{3}$$

$$(\operatorname{div} S)_z = \partial_r M + \frac{M}{r} + \partial_z C, \tag{4}$$

$$(\operatorname{div} S)_\theta = 0.$$

The $1/r$ terms are geometric, not lower-order terms that may be discarded near a shrinking cylindrical support.

Theorem 2.1 (Explicit compact meridional inversion)**.** *Let $f_r, f_z \in C_c^\infty(\Omega)$, and extend $f = f_r e_r + f_z e_z$ by zero to $\mathbb{R}^3$ in cylindrical coordinates. There is a symmetric $S \in C_c^\infty(\mathbb{R}^3; \mathrm{Sym}_3)$, axisymmetric and supported in the same open annular cylinder, with $\operatorname{div} S = f$, if and only if*

$$\int_a^b \int_c^d r f_z(r, z)\, \mathrm{d}z\, \mathrm{d}r = 0. \tag{5}$$

The construction uses no nonlocal spatial projection, but it asserts no positive-semidefinite or weighted-cone bound on S.

Proof. For necessity, integrate the axial component of $\operatorname{div} S = f$ against the cylindrical Jacobian r. Compact support and (4) give

$$\int_a^b \int_c^d r f_z\, \mathrm{d}z\, \mathrm{d}r = \int_a^b \int_c^d \left[\partial_r(rM) + r\partial_z C\right] \mathrm{d}z\, \mathrm{d}r = 0.$$

For sufficiency, define the axial source marginal and its radial primitive by

$$J(r) = \int_c^d f_z(r, z)\, \mathrm{d}z, \qquad H(r) = \frac{1}{r} \int_a^r s J(s)\, \mathrm{d}s. \tag{6}$$

Since $r > 0$ on $[a, b]$, H is smooth. The vanishing moment (5) makes H identically zero outside a compact subinterval of (a, b). Choose any $\eta \in C_c^\infty((c, d))$ normalized by $\int_c^d \eta(z)\,\mathrm{d}z = 1$, and set

$$M(r, z) = H(r)\eta(z). \tag{7}$$

The identity $(rH)' = rJ$ implies

$$\int_c^d \left(f_z - \partial_r M - \frac{M}{r} \right) \mathrm{d}z = J - H' - \frac{H}{r} = 0.$$

Consequently the axial primitive

$$C(r, z) = \int_c^z \left[f_z(r, \zeta) - \partial_r M(r, \zeta) - \frac{M(r, \zeta)}{r} \right] \mathrm{d}\zeta \tag{8}$$

vanishes both below and above the axial support; it is smooth and compactly supported in Ω. Choose any $K \in C_c^\infty(\Omega)$, including $K = 0$, and define

$$L = K + r\big(\partial_r K + \partial_z M - f_r\big). \tag{9}$$

Each component of (2) is smooth and compactly supported away from the axis. Substitution in (3) gives

$$\partial_r K + \frac{K - L}{r} + \partial_z M = \partial_r K - (\partial_r K + \partial_z M - f_r) + \partial_z M = f_r,$$

and (8) gives $\partial_r M + M/r + \partial_z C = f_z$. Thus $\operatorname{div} S = f$ as claimed. □

Remark 2.2 (Scope of the construction)**.** The theorem concerns a meridional source with $f_\theta = 0$ and support strictly separated from $r = 0$. An azimuthal source introduces a different angular-momentum compatibility condition, and extending the construction across the axis requires separate regularity conditions on the cylindrical components. The tensor in the theorem may be indefinite, and its size can deteriorate when the support shrinks. None of these omissions is harmless for the weighted H8 attachment problem.

3 Pressure-neutral gauges and a sign obstruction

The mixed and hoop degrees of freedom in Theorem 2.1 are not optional decorations. A restricted diagonal ansatz exhibits a genuine obstruction to maintaining one sign.

Proposition 3.1 (Two-diagonal gauge)**.** *Let $F \in C_c^\infty(\Omega)$, put $K = rF_r$ and $C = F + rF_r$, and take $S = Ke_r \otimes e_r + Ce_z \otimes e_z$. Then*

$$\operatorname{div} S = \nabla C. \tag{10}$$

In particular its divergence is a pressure gradient, so it cannot by itself represent a prescribed source with a nonzero solenoidal part.

Proof. The radial divergence is $K_r + K/r$. Since $K = rF_r$, $C_r = F_r + K_r = K/r + K_r$. The axial divergence is C_z, which proves (10). □

Proposition 3.2 (Restricted radial sign condition)**.** *If K and C are compactly supported in $r \in (a, b)$ at a fixed z and obey $C_r = K_r + K/r$, then*

$$\int_a^b K(r, z)/r\,\mathrm{d}r = 0. \tag{11}$$

Consequently, $K \geq 0$ forces $K \equiv 0$ for this restricted two-diagonal class.

Proof. Integrate the derivative identity between a and b; the endpoint terms of K and C vanish by compact support. Since $r > 0$, continuity and (11) imply that a nonnegative K is identically zero. □

This conclusion does not exclude a sign-changing correction around a positive background tensor or a compact construction with nonzero hoop and mixed components. Indeed, Theorem 2.1 explicitly uses both of those components. The more difficult question is whether one can choose them while retaining a prescribed strict coefficient cone with uniform shrinking-support estimates.

4 The exact seven-ray coefficient margin

The finite cone appearing in the current H8 source has seven rays. With aperture $t = 1/4$ they are

$$\begin{aligned} v_1 &= (1,0,0), \quad v_2 = (1,t,0), \quad v_3 = (1,-t,0), \\ v_4 &= (1,0,t), \quad v_5 = (1,0,-t), \\ v_6 &= (1,t,t), \quad v_7 = (1,t,-t). \end{aligned} \tag{12}$$

Their normalized positive coefficient centre is $c_i^{(0)} = 236/1229$ for $i = 3$ and $172/1229$ for every other index. The raw source covariance and its trace-free rank-one identity are established by exact finite arithmetic and in `UnforcedPatching.SevenRayCenter.lean`.

Write $(A, B, C, D) \in \mathbb{R}^4$ for the four trace-free stress coordinates. The resulting coefficient perturbations, in ray order, are

$$\delta c = \left(-\tfrac{3}{2}D,\ 2C - 4D,\ -2C - 4D,\ 2A - 8B + 4D,\ -2A + 8B + 4D,\ 8B,\ -8B\right). \tag{13}$$

Proposition 4.1 (Sharp componentwise cone estimate)**.** *For $m = \max(|A|, |B|, |C|, |D|)$ the map (13) satisfies*

$$\max_{1 \le i \le 7} |\delta c_i| \le 14m, \tag{14}$$

and the number 14 is optimal. If $\rho \ge 0$ and

$$m \le \frac{15971}{1720600}\rho, \tag{15}$$

then for every ray $c_i^{(0)}\rho + \delta c_i \ge \rho/100$.

Proof. The absolute row sums in (13) are respectively 3/2, 6, 6, 14, 14, 8, 8; hence the operator norm from $\ell^\infty(\mathbb{R}^4)$ to $\ell^\infty(\mathbb{R}^7)$ is at most 14. At $(A, B, C, D) = (1, -1, 0, 1)$, the fourth coordinate equals 14, so the constant is sharp. Since

$$\frac{172}{1229} - \frac{1}{100} = \frac{15971}{122900},$$

inequality (15) implies $|\delta c_i| \le 14m \le (15971/122900)\rho$. As $c_i^{(0)} \ge 172/1229$, the claimed floor follows. □

The seven-ray claim is a finite linear-algebra statement. It provides the exact smallness target for a potential compact correction but gives no estimate for the particular tensor delivered by Theorem 2.1. In particular, no implication from unweighted compact solvability to (15) is available.

Proposition 4.2 (The frozen-frame physical tensor map)**.** *In the local co-rotating frame* $(e_1, e_2, e_3) = (e_r, e_\theta, e_z)$ *of the seven-ray construction, let* $G = \sum_{i=1}^{7} \delta c_i v_i \otimes v_i$*. The exact physical tensor perturbation and its trace-free pressure representative are*

$$G = \begin{pmatrix} -\frac{3}{2}D & C & A \\ C & -\frac{1}{2}D & B \\ A & B & \frac{1}{2}D \end{pmatrix}, \quad H = G + \frac{D}{2}I = \begin{pmatrix} -D & C & A \\ C & 0 & B \\ A & B & D \end{pmatrix}. \tag{16}$$

In particular, if $m = \max(|A|, |B|, |C|, |D|)$*, then* $\max_{i,j} |H_{ij}| = m$ *in that frame. Conversely, the coordinates are recovered as* $A = H_{rz}$*,* $B = H_{\theta z}$*,* $C = H_{r\theta}$*, and* $D = H_{zz}$*. The image is the four-dimensional subspace of symmetric trace-free matrices with* $H_{\theta\theta} = 0$ *(hence* $H_{rr} = -H_{zz}$*); one independent trace-free diagonal mode is absent. For a different orthonormal frame with columns* E*, the physical matrix is* EHE^{T}*; its componentwise maximum need not equal* m*.*

Proof. Substitute the seven vectors (12) and their coefficient perturbations (13) into the six independent entries of G. The diagonal entries are $-3D/2, -D/2, D/2$; the off-diagonal entries in $(r\theta, rz, \theta z)$ order are (C, A, B). Thus $\operatorname{tr} G = -3D/2$, and the scalar gauge addition $DI/2$ removes its trace. The maximum-entry identity then follows directly from the displayed matrix. Conjugation gives the last claim. □

Proposition 4.3 (The exact original seven-coefficient cone)**.** *For any* $\tau \geq 0$*, all seven constraints* $|\delta c_i| \leq \tau$ *of* (13) *hold if and only if*

$$4|D| + 2|C| \leq \tau, \qquad 4|D| + 2|A - 4B| \leq \tau, \qquad 8|B| \leq \tau. \tag{17}$$

The sufficient bound $m \leq \tau/14$ *therefore need not hold for a valid coefficient perturbation: at* $(A, B, C, D) = (\tau, \tau/8, 0, 0)$*, the exact largest coefficient perturbation and* m *both equal* τ*.*

Proof. For real x, y the identity $\max(|x + y|, |x - y|) = |x| + |y|$ is exact. Apply it to the paired coefficients $(\delta c_2, \delta c_3)$, $(\delta c_4, \delta c_5)$, and $(\delta c_6, \delta c_7)$. The first pair contributes $4|D| + 2|C|$, the second $4|D| + 2|A - 4B|$, and the third $8|B|$. The remaining coefficient $|\delta c_1| = 3|D|/2$ is dominated by the first pair. Substitution verifies the equality case. □

This identifies a previously implicit assumption: the bound (15) is also a physical entrywise bound for the trace-free perturbation *in the specified cylindrical frame*. The same conclusion requires a verified change of frame when H8 rays are tilted; pointwise orthogonal conjugation alone does not control the additional derivatives of a spatially varying frame.

5 A source-matched weighted dual obstruction

The coefficient condition (15) involves four trace-free coordinates, not arbitrary physical tensor components. To test a particularly strong candidate estimate, let

$$P(r) = \begin{cases} (r-1)^2(2-r)^2, & 1 \leq r \leq 2, \\ 0, & \text{otherwise}, \end{cases} \qquad \eta(z) = \begin{cases} (1-z^2)^3, & |z| \leq 1, \\ 0, & \text{otherwise}. \end{cases} \tag{18}$$

Set $\rho(r, z) = P(r)\eta(z) \geq 0$ and choose the meridional source $f = P(r)\eta'(z)e_z$. This is source-matched in the literal sense that $S = \rho\, e_z \otimes e_z$ is a compact symmetric inverse: $\operatorname{div} S = f$. Also $\int r f_z \, \mathrm{d}r \, \mathrm{d}z = 0$, so the force compatibility condition of Theorem 2.1 holds.

Proposition 5.1 (A pressure-independent stress bound)**.** *Let $\beta = 15971/1720600$. No smooth compact symmetric tensor H and smooth scalar pressure q satisfy both*

$$\begin{aligned} \operatorname{div} H + \nabla q &= f, \\ |H_{ij}| &\le \beta\rho \quad (i,j \in \{r,\theta,z\}), \end{aligned} \tag{19}$$

where components refer to the orthonormal cylindrical frame*. The contradiction also persists for suitable smooth compact nonnegative approximations to the profiles in* (18)*.*

Proof. Construct an axisymmetric vector field with components

$$v_r = -rP\eta'', \qquad v_\theta = 0, \qquad v_z = (2P + rP')\eta'. \tag{20}$$

It is divergence-free because $r^{-1}\partial_r(rv_r)+\partial_z v_z = 0$. The zero-extended profiles give a compactly supported Lipschitz v: $P = P' = 0$ at $r = 1, 2$, and $\eta' = \eta'' = 0$ at $z = \pm 1$. Thus integration by parts against v is legitimate, by approximation if necessary, and the pressure term vanishes. With the angular factor 2π suppressed, the source pairing is

$$\begin{aligned} L &= \int_1^2 \int_{-1}^1 r f_z v_z \,\mathrm{d}z\,\mathrm{d}r \\ &= \left(\int_1^2 rP^2\,\mathrm{d}r\right)\left(\int_{-1}^1 (\eta')^2\,\mathrm{d}z\right) = \frac{1}{420}\,\frac{1024}{385} = \frac{256}{40425}. \end{aligned} \tag{21}$$

Here the middle equality follows by integrating r^2PP' by parts. The physical cylindrical gradient of v includes its hoop entry v_r/r. Its entrywise absolute sum is bounded by

$$|\nabla v|_1 \le (4P + 2r|P'|)|\eta''| + rP|\eta'''| + (3|P'| + r|P''|)|\eta'|. \tag{22}$$

The elementary exact radial integrals are $\int_1^2 P = 1/30$, $\int_1^2 P^2 = 1/630$, $\int_1^2 rP^2 = 1/420$, and $\int_1^2 P|P'| = 1/256$. On their supports, $r \le 2$, $|P''| \le 2$, $|\eta''| \le 6$ and $|\eta'''| \le 72|z| + 120|z|^3$. Direct integration gives

$$\int_{-1}^1 \eta = \frac{32}{35}, \quad \int_{-1}^1 \eta|\eta'| = 1, \quad \int_{-1}^1 \eta|z| = \frac{1}{4}, \quad \int_{-1}^1 \eta|z|^3 = \frac{1}{20}.$$

Consequently, writing $U = \int_1^2 \int_{-1}^1 r\rho|\nabla v|_1 \,\mathrm{d}z\,\mathrm{d}r$, one obtains the rigorous upper bound

$$U \le \left(\frac{4}{420} + \frac{2}{64}\right)\frac{192}{35} + \frac{2}{315}\,24 + \frac{3}{128} + \frac{4}{15} = \frac{313361}{470400}. \tag{23}$$

But (19) would imply $L \le \beta U$ by the virtual-work identity $\int f \cdot v = -\int H : \nabla v$, whereas

$$L - \beta\,\frac{313361}{470400} = \frac{443015653}{2967690880000} > 0, \tag{24}$$

a contradiction. Nonnegative convolution approximations of P and η, supported away from the axis, converge in the requisite derivative integral norms. The strict gap survives for sufficiently small smoothing parameters and proves the smooth-profile assertion. □

Corollary 5.2 (Failure of the sufficient four-field envelope)**.** *For the source f and envelope ρ in* (18)*, there is no fixed-frame four-coordinate stress $H(A, B, C, D)$ of* (16) *satisfying* $\operatorname{div} H + \nabla q = f$ *and the sufficient bound* (15)*, even if the pressure is freely chosen.*

Proof. Proposition 4.2 transfers (15) to the entrywise physical cap (19). This contradicts Proposition 5.1. □

The preceding argument excludes the *stronger sufficient* four-field estimate. It cannot alone exclude the original absolute coefficient budget $\max_i |\delta c_i| \le (15971/122900)\rho$: if one simply inserts that larger constant into the conservative gradient bound (23), then

$$\frac{15971}{122900}\frac{313361}{470400} - L = \frac{51024386641}{635933760000} > 0. \tag{25}$$

The sign of this bound is inconclusive, not evidence of existence. The exact cone, however, has a smaller sharp dual support than that componentwise comparison.

Proposition 5.3 (Sharp dual of the original absolute coefficient cone)**.** *For $\tau \ge 0$ let $\mathcal{K}_\tau$ be the four-coordinate set defined by (17). For arbitrary $a, b, c, d \in \mathbb{R}$,*

$$\sup_{(A,B,C,D)\in\mathcal{K}_\tau} (aA + bB + cC + dD) = \tau\left[\max\left(\frac{|a|+|c|}{2}, \frac{|d|}{4}\right) + \frac{|4a+b|}{8}\right]. \tag{26}$$

The upper bound in (26) is proved in Lean for the actual seven coefficient constraints.

Proof. Put $x = A - 4B$. By (17), for fixed $s = |D| \in [0, \tau/4]$, the variables x, C, B are independent and satisfy

$$|x|, |C| \le \frac{\tau - 4s}{2}, \qquad |B| \le \frac{\tau}{8}.$$

The objective becomes $ax + (4a+b)B + cC + dD$. Its maximum for this s is

$$\frac{\tau - 4s}{2}(|a| + |c|) + \frac{\tau}{8}|4a + b| + s|d|.$$

This expression is affine in s, so its maximum occurs at $s = 0$ or $s = \tau/4$, yielding (26). Each bound is attained by choosing the signs of x, C, B, D to match their respective linear coefficients. Thus the value is sharp. □

Proposition 5.4 (An exact original-budget obstruction for the manufactured source)**.** *Let f, ρ be given by (18), and set $\gamma = 15971/122900$. No locally integrable symmetric tensor of the fixed-frame form $H(A, B, C, D)$ in (16), together with a locally integrable scalar pressure q, can satisfy in the distributional sense*

$$\operatorname{div} H + \nabla q = f, \qquad |\delta c_i(A, B, C, D)| \le \gamma\rho \quad (1 \le i \le 7).$$

Proof. Use the divergence-free meridional field v from (20). Its nonzero coupling coordinates to H are

$$a(r, z) = \partial_z v_r + \partial_r v_z = -rP\eta''' + (3P' + rP'')\eta',$$
$$d(r, z) = \partial_z v_z - \partial_r v_r = (3P + 2rP')\eta''.$$

Indeed, $H : \nabla v = Aa + Dd$; the B and C components vanish because v is meridional and independent of θ. By Proposition 5.3, integration by parts and $\rho \ge 0$ would force

$$L \le \gamma S, \quad S = \int_1^2 \int_{-1}^1 rP(r)\eta(z) \max\left(|a|, \frac{|a|}{2} + \frac{|d|}{4}\right) \mathrm{d}z\, \mathrm{d}r. \tag{27}$$

There is an exact rational upper bound

$$S \le \frac{5838537817787225523484624248207}{150323855360000000000000000000000} < \frac{1}{25}. \tag{28}$$

For completeness, partition $[1, 2] \times [-1, 1]$ into 40^2 rational rectangles. On each rectangle express the separated polynomials a and d in a common-degree bivariate Bernstein basis. The

convex-hull property bounds $|a|$ and $|d|$ by the largest absolute Bernstein coefficient. Multiply $\max(a_{ij}, a_{ij}/2 + d_{ij}/4)$ by the *exact* nonnegative cell mass $\int rP\,\mathrm{d}r \int \eta\,\mathrm{d}z$, and sum the 1600 rational terms. This yields the first inequality in (28); the second is an integer comparison. The accompanying pure-rational program reproduces each coefficient and sum. Consequently

$$L - \gamma S > \frac{256}{40425} - \frac{15971}{122900 \cdot 25} = \frac{5637293}{4968232500} > 0,$$

contradicting (27). The pressure term vanishes by solenoidality of v. The zero-extended test field is compactly supported and Lipschitz, so smooth divergence-free mollifications justify the same integration-by-parts identity for locally integrable H and q. □

This obstruction concerns the *absolute correction budget* $|\delta c_i| \le \gamma\rho$ in one fixed frame. It does not rule out every nonnegative seven-ray stress: the one-sided requirement $c_i^{(0)}\rho + \delta c_i \ge \rho/100$ admits perturbations outside that sufficient absolute budget. In addition, $f = P\eta' e_z$ is a manufactured diagnostic source, *not* the authenticated H8 residual $g - \mathbb{P}\,\mathrm{div}(\rho R_*)$, where $\mathbb{P}$ denotes the divergence-free Leray projection; the source profiles and potentially rotating physical frame required to test that residual have not been provided in the checked project. No statement here excludes a more flexible source-matched H8 inverse or decides the global unforced problem.

6 The physical residual and terminal speed

Define the viscosity-one residual

$$\mathcal{N}(u,p) = u_t + (u\cdot\nabla)u - \Delta u + \nabla p. \tag{29}$$

For smooth (u,p) and (w,q), direct expansion yields

$$\mathcal{N}(u+w, p+q) = \mathcal{N}(u,p) + \mathcal{N}(w,q) + (u\cdot\nabla)w + (w\cdot\nabla)u. \tag{30}$$

Both nonlinear cross terms are essential. Equation (30), with explicit differentiability assumptions in the upstream physical definitions, is checked in Lean by the declaration `UnforcedPatching.navierStokesResidual_add`.

Proposition 6.1 (Conditional terminal retention)**.** *Suppose that u has arbitrarily large speed at points with $t \uparrow 1$. If a correction w satisfies*

$$\|w(t,x)\| \le \alpha\|u(t,x)\| + C_0, \qquad 0 \le \alpha < 1, \quad C_0 \ge 0, \tag{31}$$

uniformly on $0 < t < 1$ and $x \in \mathbb{T}^3$, then $u + w$ also has arbitrarily large speed arbitrarily close to time one.

Proof. The reverse triangle inequality and (31) give

$$\|u+w\| \ge \|u\| - \|w\| \ge (1-\alpha)\|u\| - C_0.$$

Given any target speed $M > 0$ and any terminal neighbourhood, choose a point therein where $\|u\| > (M + C_0)/(1-\alpha)$. Then $\|u+w\| > M$ at that point. □

This last implication is kernel-checked by the Lean theorem

`UnforcedPatching.speedUnbounded_add_subleading`

It is deliberately conditional: no field u meeting the singularity premise, and no actual PDE correction w meeting (31), has been supplied by these lemmas.

7 What the formal development proves

The Lean project is pinned to Lean 4.34.0-rc2 and a fixed Mathlib revision. The checked physical interface states smoothness, three coordinate periods, zero forcing, nonzero initial data, finite fundamental-cell kinetic energy and endpoint speed divergence. Its existential proposition is named `UnforcedPatching.unforcedTargetStatement`; there is no proof term inhabiting that proposition.

The precise correspondence between the paper and the compiled `UnforcedPatching` declarations is as follows. In the second column, "missing" refers to mathematical constructions, not to unresolved Lean goals inside the named theorems.

Lean module	Checked scope and the missing construction
`Definitions`	Physical candidate predicates; no candidate witness.
`Residual`, `Correction`	Full residual addition and conditional cancellation; no constructed correcting field.
`Periodicity`, `PressureGauge`	Sum and time-only gauge identities; no periodic singular solution.
`ConeBudget`, `ConeOperator`, `SevenRayCenter`	Exact rational centre and sharp factor 14; no source-matched weighted positive stress.
`FourFieldPhysical`	Exact nine-entry ray-to-cylindrical tensor map modulo an isotropic gauge; entrywise four-coordinate bound in the frozen frame, not a spatially rotated PDE bridge.
`ExactCoefficientCone`, `SharpCoefficientDual`	Exact seven-coefficient polytope and universal sharp dual upper estimate; no authentic H8 source or spatial inverse.
`CompactGauge`	One-variable gauge derivative; no compact tensor support theorem.
`MeridionalInverse`	Actual axial and radial interval-integral FTC identities, geometric hoop cancellation, and pointwise radial/axial cylindrical inverse; no Cartesian tensor-divergence bridge or compact support theorem.
`Retention`	Subleading corrections preserve terminal speed; neither required field constructed.

In particular, the residual-addition, correction-closure, seven-ray-floor, and terminal-retention declarations are all implications with explicit hypotheses. Their conjunction does not yield an inhabitant of `unforcedTargetStatement`. The printed axioms for the selected smoke-tested declarations are the standard logical axioms `propext`, `Classical.choice`, and `Quot.sound`.

In particular, `meridional_correction_divergence` checks both components of (3)–(4) for the hoop stress (9) and the actual interval-integral axial primitive (8), under stated local continuity and integrability assumptions. The separate declarations `radialMoment_relation` and `mixedCarrier_radial_balance` verify the antiderivative identity in (6) and its product carrier (7). The further theorem `sourceMatched_axialStress` checks the pointwise cylindrical source match $S_{zz} = \rho$, $f_z = \partial_z \rho$ used in Proposition 5.1. These results compile without proof admissions; selected axiom reports list only Lean's standard logical axioms. They remain pointwise statements about *defined cylindrical expressions*: the Cartesian tensor-divergence equivalence, global support and regularity construction, moment compatibility and weighted H8 estimate are not yet formalized in Lean. Thus the full restricted Theorem 2.1 is *partially*, not fully, formalized.

An independent exact-arithmetic test evaluates a nontrivial annular polynomial instance of the radial compatibility identity. For the uncompensated purely radial version, the curl obstruction at $(r, z) = (5/4, 1/2)$ equals $-5103/32768$ exactly; the test explicitly rejects

the pressure-neutral interpretation in this case. The polynomial is used only for interior differentiation, not as an infinitely smooth zero-extended cutoff.

A second independent exact-rational test reproduces (21)–(24) from polynomial integrals and checks the complete cylindrical divergence of the dual test field. An additional exact reconstruction and Lean theorem prove the fixed-frame four-field map (16), its inverse coordinate recovery and its entrywise consequence; the test (25) records the failure of this *conservative gradient method* at the original coefficient budget. The sharper polytope method proves Proposition 5.4: a 40×40 Bernstein partition yields the exact rational upper bound (28). No executable test substitutes for the authentic H8 source and rotated-frame PDE estimates. The equivalence (17) is independently checked on exact rational samples and kernel-checked by `original_seven_coefficient_cone_exact`. The additional kernel theorem `original_cone_dual_upper` establishes the finite-dimensional upper bound in (26); exact rational maximizing witnesses independently verify sharpness. The Bernstein integral itself is a rational Python certificate, not a Lean theorem.

The source code, machine-readable G0–G8 ledger, smoke tests and reproduction commands accompany the public integration branch [6]. A local pinned build has passed fourteen smoke modules; at the time of preparing this edition, the corresponding GitHub-hosted Actions checks had not furnished a successful receipt for the latest commit. Local compilation does not replace hosted CI, semantic theorem auditing or independent scrutiny of the paper proof.

8 Remaining analytic requirements

The compact meridional inverse of Theorem 2.1 is a properly restricted solvability result. It does not resolve the weighted attachment problem because the desired source may have an azimuthal component, a noncompact Leray-projected tail, nontrivial pressure coupling, or incompatible moments. Even when compact inversion is possible, the output may leave the required seven-ray positive cone, especially as the support shrinks. The strict dual obstruction in Proposition 5.1, together with Corollary 5.2, excludes the sufficient four-coordinate estimate; Proposition 5.4 also excludes the original absolute coefficient correction budget for that same manufactured source. The next mathematical question is whether the *authentic* source, perhaps using a tilted frame or a more flexible coefficient allocation, admits a weighted inverse satisfying the one-sided positive H8 ray conditions (or the absolute inequalities (17)) and all spatial derivative bounds. The sharp cone description enables a direct weighted dual optimization against the authentic source, without replacing those inequalities by the factor-14 sufficient estimate.

A complete unforced singularity construction must still provide a single authenticated local profile, source-matched smooth shell data, a scale-sharp weighted symmetric inverse respecting the coefficient margin, a periodic nonlinear correction with a proved right-inverse/contraction scheme, the full physical zero-residual identity, and terminal retention for that actual correction. Each claim requires a further source-to-Lean semantic audit and a compiled theorem using the genuine PDE definitions.

Thus the present mathematical result is an exact compact stress construction and several independently checked necessary or conditional ingredients. The existence or nonexistence of a finite-time singular solution to (1) with smooth nonzero periodic initial data is *not decided by this manuscript.*

References

[1] C. Fefferman, *Existence and Smoothness of the Navier–Stokes Equation*, Clay Mathematics Institute Millennium Prize Problem statement, https://www.claymath.org/wp-content/uploads/2022/06/navierstokes.pdf.

[2] OpenAI, *Finite time blowup for Navier–Stokes* (2026), and the associated formalization repository, https://github.com/openai/NavierStokesAndEuler.

[3] X. Zhao, *Compactly Supported Finite-Energy Stationary Solutions of the Three-Dimensional Navier–Stokes Equations* (2026), arXiv:2608.17383, https://arxiv.org/abs/2608.17383.

[4] J. Leray, *Sur le mouvement d'un liquide visqueux emplissant l'espace*, Acta Mathematica **63** (1934), 193–248.

[5] L. Caffarelli, R. Kohn, and L. Nirenberg, *Partial regularity of suitable weak solutions of the Navier–Stokes equations*, Communications on Pure and Applied Mathematics **35** (1982), 771–831.

[6] T. Newton, *Unforced Navier–Stokes Lean integration branch*, https://github.com/tonynewton79-web/unforced-navier-stokes-lean/pull/3.